

SUBLINEAR TIME LOW-RANK APPROXIMATION OF POSITIVE SEMIDEFINITE MATRICES

Cameron Musco (MIT) and David P. Woodruff (CMU)

Our Contributions:

Our Contributions:

- A near optimal low-rank approximation for any positive semidefinite (PSD) matrix can be computed in sublinear time (i.e. without reading the full matrix).

Our Contributions:

- A near optimal low-rank approximation for any positive semidefinite (PSD) matrix can be computed in sublinear time (i.e. without reading the full matrix).
- **Concrete:** Significantly improves on previous, roughly linear time approaches for general matrices, and bypasses a trivial linear time lower bound for general matrices.

Our Contributions:


- A near optimal low-rank approximation for any positive semidefinite (PSD) matrix can be computed **in sublinear time** (i.e. without reading the full matrix).
- **Concrete:** Significantly improves on previous, roughly linear time approaches for general matrices, and bypasses a trivial linear time lower bound for general matrices.
- **High Level:** Demonstrates that PSD structure can be exploited in a much stronger way than previously known for low-rank approximation. Opens the possibility of further advances in algorithms for PSD matrices.

LOW-RANK MATRIX APPROXIMATION

Low-rank approximation is one of the most widely used methods for general matrix and data compression.


LOW-RANK MATRIX APPROXIMATION

Low-rank approximation is one of the most widely used methods for general matrix and data compression.



LOW-RANK MATRIX APPROXIMATION


Low-rank approximation is one of the most widely used methods for general matrix and data compression.



- Closely related to principal component analysis, spectral embedding/clustering, and low-rank matrix completion.

LOW-RANK MATRIX APPROXIMATION

Low-rank approximation is one of the most widely used methods for general matrix and data compression.




- Closely related to principal component analysis, spectral embedding/clustering, and low-rank matrix completion.

Important Special Case: A is positive semidefinite (PSD). I.e.

$$\mathbf{x}^T \mathbf{A} \mathbf{x} \geq 0, \forall \mathbf{x} \in \mathbb{R}^n$$

LOW-RANK MATRIX APPROXIMATION

Low-rank approximation is one of the most widely used methods for general matrix and data compression.



- Closely related to principal component analysis, spectral embedding/clustering, and low-rank matrix completion.

Important Special Case: A is positive semidefinite (PSD). I.e.


$$\mathbf{x}^T \mathbf{A} \mathbf{x} \geq 0, \forall \mathbf{x} \in \mathbb{R}^n$$

- Includes graph Laplacians, Gram matrices and kernel matrices, covariance matrices, Hessians for convex functions.

An optimal low-rank approximation can be computed via the singular value decomposition (SVD).


OPTIMAL LOW-RANK APPROXIMATION

An optimal low-rank approximation can be computed via the singular value decomposition (SVD).




OPTIMAL LOW-RANK APPROXIMATION

An optimal low-rank approximation can be computed via the singular value decomposition (SVD).



OPTIMAL LOW-RANK APPROXIMATION


An optimal low-rank approximation can be computed via the singular value decomposition (SVD).



$$\mathbf{A}_k = \arg \min_{\mathbf{B}: \text{rank}(\mathbf{B})=k} \|\mathbf{A} - \mathbf{B}\|_F$$

OPTIMAL LOW-RANK APPROXIMATION


An optimal low-rank approximation can be computed via the singular value decomposition (SVD).



$$\mathbf{A}_k = \arg \min_{\mathbf{B}: \text{rank}(\mathbf{B})=k} \|\mathbf{A} - \mathbf{B}\|_F = \sqrt{\sum_{i,j} (\mathbf{A}_{ij} - \mathbf{B}_{ij})^2}$$

OPTIMAL LOW-RANK APPROXIMATION

An optimal low-rank approximation can be computed via the singular value decomposition (SVD).



$$\mathbf{A}_k = \arg \min_{\mathbf{B}: \text{rank}(\mathbf{B})=k} \|\mathbf{A} - \mathbf{B}\|_F = \sqrt{\sum_{i,j} (\mathbf{A}_{ij} - \mathbf{B}_{ij})^2}$$

- Unfortunately, computing the SVD takes $O(nd^2)$ time.

Recent work on matrix sketching gives state-of-the-art runtimes

Recent work on matrix sketching gives state-of-the-art runtimes

Theorem (Clarkson, Woodruff '13)

There is an algorithm which in $O(\text{nnz}(\mathbf{A}) + n \cdot \text{poly}(k, 1/\epsilon))$ time outputs $\mathbf{N} \in \mathbb{R}^{n \times k}$, $\mathbf{M} \in \mathbb{R}^{d \times k}$ satisfying with prob. 99/100:

$$\|\mathbf{A} - \mathbf{NM}^T\|_F \leq (1 + \epsilon)\|\mathbf{A} - \mathbf{A}_k\|_F.$$

Recent work on matrix sketching gives state-of-the-art runtimes

Theorem (Clarkson, Woodruff '13)

There is an algorithm which in $O(\text{nnz}(\mathbf{A}) + n \cdot \text{poly}(k, 1/\epsilon))$ time outputs $\mathbf{N} \in \mathbb{R}^{n \times k}$, $\mathbf{M} \in \mathbb{R}^{d \times k}$ satisfying with prob. 99/100:

$$\|\mathbf{A} - \mathbf{NM}^T\|_F \leq (1 + \epsilon) \|\mathbf{A} - \mathbf{A}_k\|_F.$$

- When $k, 1/\epsilon$ are not too large, runtime is linear in input size.

Recent work on matrix sketching gives state-of-the-art runtimes

Theorem (Clarkson, Woodruff '13)

There is an algorithm which in $O(\text{nnz}(\mathbf{A}) + n \cdot \text{poly}(k, 1/\epsilon))$ time outputs $\mathbf{N} \in \mathbb{R}^{n \times k}$, $\mathbf{M} \in \mathbb{R}^{d \times k}$ satisfying with prob. 99/100:

$$\|\mathbf{A} - \mathbf{NM}^T\|_F \leq (1 + \epsilon)\|\mathbf{A} - \mathbf{A}_k\|_F.$$

- When $k, 1/\epsilon$ are not too large, runtime is **linear in input size**.
- Best known runtime for both general and PSD matrices.

Theorem (Main Result – Musco, Woodruff '17)

*There is an algorithm running in $\tilde{O}\left(\frac{nk^2}{\epsilon^4}\right)$ time which, given *PSD* \mathbf{A} , outputs $\mathbf{N}, \mathbf{M} \in \mathbb{R}^{n \times k}$ satisfying with probability 99/100:*

$$\|\mathbf{A} - \mathbf{NM}^T\|_F \leq (1 + \epsilon)\|\mathbf{A} - \mathbf{A}_k\|_F.$$

Theorem (Main Result – Musco, Woodruff '17)

There is an algorithm running in $\tilde{O}\left(\frac{nk^2}{\epsilon^4}\right)$ time which, given *PSD* \mathbf{A} , outputs $\mathbf{N}, \mathbf{M} \in \mathbb{R}^{n \times k}$ satisfying with probability 99/100:

$$\|\mathbf{A} - \mathbf{NM}^T\|_F \leq (1 + \epsilon)\|\mathbf{A} - \mathbf{A}_k\|_F.$$

- Compare to CW'13 which takes $O(\text{nnz}(\mathbf{A})) + n \cdot \text{poly}(k, 1/\epsilon)$.

Theorem (Main Result – Musco, Woodruff '17)

There is an algorithm running in $\tilde{O}\left(\frac{nk^2}{\epsilon^4}\right)$ time which, given *PSD* \mathbf{A} , outputs $\mathbf{N}, \mathbf{M} \in \mathbb{R}^{n \times k}$ satisfying with probability 99/100:

$$\|\mathbf{A} - \mathbf{NM}^T\|_F \leq (1 + \epsilon)\|\mathbf{A} - \mathbf{A}_k\|_F.$$

- Compare to CW'13 which takes $O(\text{nnz}(\mathbf{A})) + n \cdot \text{poly}(k, 1/\epsilon)$.

Theorem (Main Result – Musco, Woodruff '17)

There is an algorithm running in $\tilde{O}\left(\frac{nk^2}{\epsilon^4}\right)$ time which, given PSD \mathbf{A} , outputs $\mathbf{N}, \mathbf{M} \in \mathbb{R}^{n \times k}$ satisfying with probability 99/100:

$$\|\mathbf{A} - \mathbf{NM}^T\|_F \leq (1 + \epsilon)\|\mathbf{A} - \mathbf{A}_k\|_F.$$

- Compare to CW'13 which takes $O(\text{nnz}(\mathbf{A})) + n \cdot \text{poly}(k, 1/\epsilon)$.
- If $k, 1/\epsilon$ are not too large compared to $\text{nnz}(\mathbf{A})$, our runtime is significantly sublinear in the size of \mathbf{A} .

LOWER BOUND FOR GENERAL MATRICES

For general matrices, $\Omega(\text{nnz}(\mathbf{A}))$ time is required.


For general matrices, $\Omega(\text{nnz}(\mathbf{A}))$ time is required.

- Randomly place a single entry which dominates \mathbf{A} 's Frobenius norm.

LOWER BOUND FOR GENERAL MATRICES

For general matrices, $\Omega(\text{nnz}(\mathbf{A}))$ time is required.


- Randomly place a single entry which dominates \mathbf{A} 's Frobenius norm.



LOWER BOUND FOR GENERAL MATRICES

For general matrices, $\Omega(\text{nnz}(\mathbf{A}))$ time is required.


- Randomly place a single entry which dominates \mathbf{A} 's Frobenius norm.
- Finding it with constant probability requires reading at least a constant fraction of the non-zero entries in \mathbf{A} .



LOWER BOUND FOR GENERAL MATRICES

For general matrices, $\Omega(\text{nnz}(\mathbf{A}))$ time is required.

- Randomly place a single entry which dominates \mathbf{A} 's Frobenius norm.
- Finding it with constant probability requires reading at least a constant fraction of the non-zero entries in \mathbf{A} .




- Lower bound holds for any approximation factor and even rules out $o(\text{nnz}(\mathbf{A}))$ time for weaker guarantees.

LOWER BOUND FOR GENERAL MATRICES

For general matrices, $\Omega(\text{nnz}(\mathbf{A}))$ time is required.

- Randomly place a single entry which dominates \mathbf{A} 's Frobenius norm.
- Finding it with constant probability requires reading at least a constant fraction of the non-zero entries in \mathbf{A} .




- Lower bound holds for any approximation factor and even rules out $o(\text{nnz}(\mathbf{A}))$ time for weaker guarantees.

$$\|\mathbf{A} - \mathbf{NM}^T\|_F \leq (1 + \epsilon)\|\mathbf{A} - \mathbf{A}_k\|_F$$

LOWER BOUND FOR GENERAL MATRICES

For general matrices, $\Omega(\text{nnz}(\mathbf{A}))$ time is required.

- Randomly place a single entry which dominates \mathbf{A} 's Frobenius norm.
- Finding it with constant probability requires reading at least a constant fraction of the non-zero entries in \mathbf{A} .



- Lower bound holds for any approximation factor and even rules out $o(\text{nnz}(\mathbf{A}))$ time for weaker guarantees.


$$\|\mathbf{A} - \mathbf{NM}^T\|_F \leq \|\mathbf{A} - \mathbf{A}_k\|_F + \epsilon \|\mathbf{A}\|_F.$$

WHAT ABOUT FOR PSD MATRICES?

Observation: For PSD \mathbf{A} , we have for any entry \mathbf{a}_{ij} :

$$\mathbf{a}_{ij} \leq \max(\mathbf{a}_{ii}, \mathbf{a}_{jj})$$

since otherwise $(\mathbf{e}_i - \mathbf{e}_j)^T \mathbf{A} (\mathbf{e}_i - \mathbf{e}_j) < 0$.




WHAT ABOUT FOR PSD MATRICES?

Observation: For PSD \mathbf{A} , we have for any entry \mathbf{a}_{ij} :

$$\mathbf{a}_{ij} \leq \max(\mathbf{a}_{ii}, \mathbf{a}_{jj})$$

since otherwise $(\mathbf{e}_i - \mathbf{e}_j)^T \mathbf{A} (\mathbf{e}_i - \mathbf{e}_j) < 0$.

- So we can find any 'hidden' heavy entry by looking at its corresponding diagonal entries.




WHAT ABOUT FOR PSD MATRICES?

Observation: For PSD \mathbf{A} , we have for any entry \mathbf{a}_{ij} :

$$\mathbf{a}_{ij} \leq \max(\mathbf{a}_{ii}, \mathbf{a}_{jj})$$

since otherwise $(\mathbf{e}_i - \mathbf{e}_j)^T \mathbf{A} (\mathbf{e}_i - \mathbf{e}_j) < 0$.

- So we can find any 'hidden' heavy entry by looking at its corresponding diagonal entries.



Question: How can we exploit additional structure arising from positive semidefiniteness to achieve sublinear runtime?

EVERY PSD MATRIX IS A GRAM MATRIX

Very Simple Fact: Every PSD matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ can be written as $\mathbf{B}^T \mathbf{B}$ for some $\mathbf{B} \in \mathbb{R}^{n \times n}$.

EVERY PSD MATRIX IS A GRAM MATRIX

Very Simple Fact: Every PSD matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ can be written as $\mathbf{B}^T \mathbf{B}$ for some $\mathbf{B} \in \mathbb{R}^{n \times n}$.

- \mathbf{B} can be any matrix square root of \mathbf{A} , e.g. if we let $\mathbf{V}\mathbf{\Sigma}\mathbf{V}^T$ be the eigendecomposition of \mathbf{A} , we can set $\mathbf{B} = \mathbf{\Sigma}^{1/2}\mathbf{V}^T$.

EVERY PSD MATRIX IS A GRAM MATRIX

Very Simple Fact: Every PSD matrix $\mathbf{A} \in \mathbb{R}^{n \times n}$ can be written as $\mathbf{B}^T \mathbf{B}$ for some $\mathbf{B} \in \mathbb{R}^{n \times n}$.

- \mathbf{B} can be any matrix square root of \mathbf{A} , e.g. if we let $\mathbf{V}\mathbf{\Sigma}\mathbf{V}^T$ be the eigendecomposition of \mathbf{A} , we can set $\mathbf{B} = \mathbf{\Sigma}^{1/2}\mathbf{V}^T$.
- Letting $\mathbf{b}_1, \dots, \mathbf{b}_n$ be the columns of \mathbf{B} , the entries of \mathbf{A} contain every pairwise dot product $a_{ij} = \mathbf{b}_i^T \mathbf{b}_j$.

$$\begin{matrix} \text{ } & \mathbf{b}_i^T \\ \mathbf{B}^T & \end{matrix} \quad \begin{matrix} & & \\ & \mathbf{B} & \mathbf{b}_j \\ & & \end{matrix} = \begin{matrix} & & \mathbf{a}_{ij} \\ & \mathbf{A} & \\ & & \end{matrix}$$

EVERY PSD MATRIX IS A GRAM MATRIX

The fact that \mathbf{A} is a Gram matrix places a variety of geometric constraints on its entries.

EVERY PSD MATRIX IS A GRAM MATRIX

The fact that \mathbf{A} is a Gram matrix places a variety of **geometric constraints** on its entries.

- The heavy diagonal observation is just one example. By Cauchy-Schwarz:

$$\mathbf{a}_{ij} = \mathbf{b}_i^T \mathbf{b}_j \leq \sqrt{(\mathbf{b}_i^T \mathbf{b}_i) \cdot (\mathbf{b}_j^T \mathbf{b}_j)} = \sqrt{\mathbf{a}_{ii} \cdot \mathbf{a}_{jj}} \leq \max(\mathbf{a}_{ii}, \mathbf{a}_{jj}).$$

EVERY PSD MATRIX IS A GRAM MATRIX

The fact that \mathbf{A} is a Gram matrix places a variety of **geometric constraints** on its entries.

- The heavy diagonal observation is just one example. By Cauchy-Schwarz:

$$\mathbf{a}_{ij} = \mathbf{b}_i^T \mathbf{b}_j \leq \sqrt{(\mathbf{b}_i^T \mathbf{b}_i) \cdot (\mathbf{b}_j^T \mathbf{b}_j)} = \sqrt{\mathbf{a}_{ii} \cdot \mathbf{a}_{jj}} \leq \max(\mathbf{a}_{ii}, \mathbf{a}_{jj}).$$

Another View: \mathbf{A} contains a lot of information about the column span of \mathbf{B} in a very compressed form – with every pairwise dot product stored as \mathbf{a}_{ij} .

Question: Can we compute a low-rank approximation of \mathbf{B} using $o(n^2)$ column dot products? I.e. $o(n^2)$ accesses to \mathbf{A} ?

Question: Can we compute a low-rank approximation of \mathbf{B} using $o(n^2)$ column dot products? I.e. $o(n^2)$ accesses to \mathbf{A} ?

Why? \mathbf{B} has the same (right) singular vectors as \mathbf{A} , and its singular values are closely related: $\sigma_i(\mathbf{B}) = \sqrt{\sigma_i(\mathbf{A})}$.

Question: Can we compute a low-rank approximation of \mathbf{B} using $o(n^2)$ column dot products? I.e. $o(n^2)$ accesses to \mathbf{A} ?

Why? \mathbf{B} has the same (right) singular vectors as \mathbf{A} , and its singular values are closely related: $\sigma_i(\mathbf{B}) = \sqrt{\sigma_i(\mathbf{A})}$.

- So the top k singular vectors are the same for the two matrices. An **optimal** low-rank approximation for \mathbf{B} thus gives an optimal low-rank approximation for \mathbf{A} .

Question: Can we compute a low-rank approximation of \mathbf{B} using $o(n^2)$ column dot products? I.e. $o(n^2)$ accesses to \mathbf{A} ?

Why? \mathbf{B} has the same (right) singular vectors as \mathbf{A} , and its singular values are closely related: $\sigma_i(\mathbf{B}) = \sqrt{\sigma_i(\mathbf{A})}$.

- So the top k singular vectors are the same for the two matrices. An **optimal** low-rank approximation for \mathbf{B} thus gives an optimal low-rank approximation for \mathbf{A} .
- Things will be messier once we introduce approximation, but **this simple idea will lead to a sublinear time algorithm for \mathbf{A} .**

Theorem (Deshpande, Vempala '06)

For any $\mathbf{B} \in \mathbb{R}^{n \times n}$, there exists a subset of $\tilde{O}(k^2/\epsilon)$ columns whose span contains $\mathbf{Z} \in \mathbb{R}^{n \times k}$ satisfying:

$$\|\mathbf{B} - \mathbf{Z}\mathbf{Z}^T\mathbf{B}\|_F \leq (1 + \epsilon)\|\mathbf{B} - \mathbf{B}_k\|_F$$

Theorem (Deshpande, Vempala '06)

For any $\mathbf{B} \in \mathbb{R}^{n \times n}$, there exists a subset of $\tilde{O}(k^2/\epsilon)$ columns whose span contains $\mathbf{Z} \in \mathbb{R}^{n \times k}$ satisfying:

$$\|\mathbf{B} - \mathbf{Z}\mathbf{Z}^T\mathbf{B}\|_F \leq (1 + \epsilon)\|\mathbf{B} - \mathbf{B}_k\|_F$$

Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}}\mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}}\mathbf{b}_i\|^2}$.

Theorem (Deshpande, Vempala '06)

For any $\mathbf{B} \in \mathbb{R}^{n \times n}$, there exists a subset of $\tilde{O}(k^2/\epsilon)$ columns whose span contains $\mathbf{Z} \in \mathbb{R}^{n \times k}$ satisfying:

$$\|\mathbf{B} - \mathbf{Z}\mathbf{Z}^T\mathbf{B}\|_F \leq (1 + \epsilon)\|\mathbf{B} - \mathbf{B}_k\|_F$$

Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}}\mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}}\mathbf{b}_i\|^2}$.



Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}$.





Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}}\mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}}\mathbf{b}_i\|^2} = \frac{\|\mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i\|^2} = \frac{a_{ii}}{\text{tr}(\mathbf{A})}$.





Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}}\mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}}\mathbf{b}_i\|^2} = \frac{\|\mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i\|^2} = \frac{\mathbf{a}_{ii}}{\text{tr}(\mathbf{A})}$.





Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}}\mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}}\mathbf{b}_i\|^2} = \frac{\|\mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i\|^2} = \frac{a_{ii}}{\text{tr}(\mathbf{A})}$.





Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}$.





Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}$.





Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}}\mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}}\mathbf{b}_i\|^2}$.





Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}$.





Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}$.





Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}$.





Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}$.





Adaptive Sampling

Initially, start with an empty column subset $\mathcal{S} := \{\}$.

For $t = 1, \dots, \tilde{O}(k^2/\epsilon)$

Let $\mathbf{P}_{\mathcal{S}}$ be the projection onto the columns in \mathcal{S} .

Add \mathbf{b}_i to \mathcal{S} with probability $\frac{\|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}{\sum_{i=1}^n \|\mathbf{b}_i - \mathbf{P}_{\mathcal{S}} \mathbf{b}_i\|^2}$.



Theorem (Factor Matrix Low-Rank Approximation)

There is an algorithm using $\tilde{O}(nk^2/\epsilon)$ accesses to $\mathbf{A} = \mathbf{B}^T \mathbf{B}$ which computes $\mathbf{Z} \in \mathbb{R}^{n \times k}$ satisfying with probability 99/100:

$$\|\mathbf{B} - \mathbf{Z}\mathbf{Z}^T \mathbf{B}\|_F \leq (1 + \epsilon) \|\mathbf{B} - \mathbf{B}_k\|_F.$$

Theorem (Factor Matrix Low-Rank Approximation)

There is an algorithm using $\tilde{O}(nk^2/\epsilon)$ accesses to $\mathbf{A} = \mathbf{B}^T \mathbf{B}$ which computes $\mathbf{Z} \in \mathbb{R}^{n \times k}$ satisfying with probability 99/100:

$$\|\mathbf{B} - \mathbf{Z}\mathbf{Z}^T \mathbf{B}\|_F \leq (1 + \epsilon) \|\mathbf{B} - \mathbf{B}_k\|_F.$$

- How does this translate to low-rank approximation of \mathbf{A} itself?

Lemma

If $\|\mathbf{B} - \mathbf{Z}\mathbf{Z}^T\mathbf{B}\|_F^2 \leq \left(1 + \frac{\epsilon^{3/2}}{\sqrt{n}}\right) \|\mathbf{B} - \mathbf{B}_k\|_F^2$, then for $\mathbf{A} = \mathbf{B}^T\mathbf{B}$:

$$\|\mathbf{A} - \mathbf{B}^T\mathbf{Z}\mathbf{Z}^T\mathbf{B}\|_F^2 \leq (1 + \epsilon)\|\mathbf{A} - \mathbf{A}_k\|_F^2.$$

Lemma

If $\|\mathbf{B} - \mathbf{Z}\mathbf{Z}^T\mathbf{B}\|_F^2 \leq \left(1 + \frac{\epsilon^{3/2}}{\sqrt{n}}\right) \|\mathbf{B} - \mathbf{B}_k\|_F^2$, then for $\mathbf{A} = \mathbf{B}^T\mathbf{B}$:

$$\|\mathbf{A} - \mathbf{B}^T\mathbf{Z}\mathbf{Z}^T\mathbf{B}\|_F^2 \leq (1 + \epsilon)\|\mathbf{A} - \mathbf{A}_k\|_F^2.$$

Lemma

If $\|\mathbf{B} - \mathbf{Z}\mathbf{Z}^T\mathbf{B}\|_F^2 \leq \left(1 + \frac{\epsilon^{3/2}}{\sqrt{n}}\right) \|\mathbf{B} - \mathbf{B}_k\|_F^2$, then for $\mathbf{A} = \mathbf{B}^T\mathbf{B}$:

$$\|\mathbf{A} - \mathbf{A}\mathbf{S}\mathbf{C}\mathbf{S}^T\mathbf{A}^T\|_F^2 \leq (1 + \epsilon)\|\mathbf{A} - \mathbf{A}_k\|_F^2.$$

Lemma

If $\|\mathbf{B} - \mathbf{Z}\mathbf{Z}^T\mathbf{B}\|_F^2 \leq \left(1 + \frac{\epsilon^{3/2}}{\sqrt{n}}\right) \|\mathbf{B} - \mathbf{B}_k\|_F^2$, then for $\mathbf{A} = \mathbf{B}^T\mathbf{B}$:

$$\|\mathbf{A} - \mathbf{A}\mathbf{S}\mathbf{C}\mathbf{S}^T\mathbf{A}^T\|_F^2 \leq (1 + \epsilon)\|\mathbf{A} - \mathbf{A}_k\|_F^2.$$

- This gives a low-rank approximation algorithm which accesses just $\tilde{O}\left(\frac{nk^2}{\epsilon^{3/2}/\sqrt{n}}\right) = n^{3/2} \cdot \text{poly}(k, 1/\epsilon)$ entries of \mathbf{A} .

Lemma

If $\|\mathbf{B} - \mathbf{Z}\mathbf{Z}^T\mathbf{B}\|_F^2 \leq \left(1 + \frac{\epsilon^{3/2}}{\sqrt{n}}\right) \|\mathbf{B} - \mathbf{B}_k\|_F^2$, then for $\mathbf{A} = \mathbf{B}^T\mathbf{B}$:

$$\|\mathbf{A} - \mathbf{A}\mathbf{S}\mathbf{C}\mathbf{S}^T\mathbf{A}^T\|_F^2 \leq (1 + \epsilon)\|\mathbf{A} - \mathbf{A}_k\|_F^2.$$


- This gives a low-rank approximation algorithm which accesses just $\tilde{O}\left(\frac{nk^2}{\epsilon^{3/2}/\sqrt{n}}\right) = n^{3/2} \cdot \text{poly}(k, 1/\epsilon)$ entries of \mathbf{A} .
- Our best algorithm accesses just $\tilde{O}\left(\frac{nk}{\epsilon^{2.5}}\right)$ entries of \mathbf{A} and runs in $\tilde{O}\left(\frac{nk^2}{\epsilon^4}\right)$ time.

Recall that our algorithm accesses the diagonal of \mathbf{A} along with $\tilde{O}(k^2/\epsilon)$ columns.

Recall that our algorithm accesses the diagonal of \mathbf{A} along with $\tilde{O}(k^2\sqrt{n})$ columns.


LIMITATIONS OF COLUMN SAMPLING

Recall that our algorithm accesses the diagonal of **A** along with $\tilde{O}(k^2\sqrt{n})$ columns.



LIMITATIONS OF COLUMN SAMPLING

Recall that our algorithm accesses the diagonal of \mathbf{A} along with $\tilde{O}(k^2\sqrt{n})$ columns.



- If we take fewer columns, we can miss a $\sqrt{n} \times \sqrt{n}$ block which contains a constant fraction of \mathbf{A} 's Frobenius norm.

Solution: Sample both rows and columns of **A**.

Solution: Sample both rows and columns of \mathbf{A} .

- Instead of adaptive sampling we use **ridge leverage scores**, which can also be computed using an iterative sampling scheme making $\tilde{O}(nk)$ accesses to \mathbf{A} (Musco, Musco '17).

Solution: Sample both rows and columns of \mathbf{A} .

- Instead of adaptive sampling we use **ridge leverage scores**, which can also be computed using an iterative sampling scheme making $\tilde{O}(nk)$ accesses to \mathbf{A} (Musco, Musco '17).
- Same intuition – select a diverse set of columns which span a near-optimal low-rank approximation of the matrix.


Solution: Sample both rows and columns of \mathbf{A} .

- Instead of adaptive sampling we use **ridge leverage scores**, which can also be computed using an iterative sampling scheme making $\tilde{O}(nk)$ accesses to \mathbf{A} (Musco, Musco '17).
- Same intuition – select a diverse set of columns which span a near-optimal low-rank approximation of the matrix.
- Sample \mathbf{AS} is a **projection-cost-preserving sketch** for \mathbf{A} [Cohen et al '15,'17]. For any rank- k projection \mathbf{P} ,

$$\|\mathbf{AS} - \mathbf{PAS}\|_F^2 = (1 \pm \epsilon) \|\mathbf{A} - \mathbf{PA}\|_F^2.$$

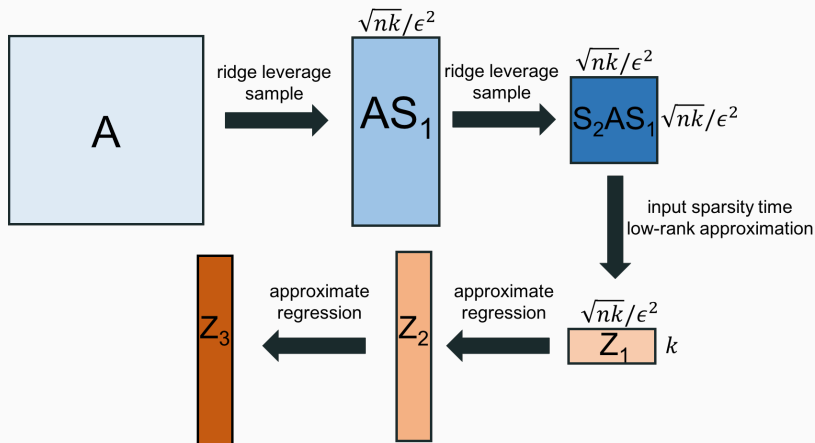
Recover low-rank approximation using two-sided sampling and projection-cost-preserving sketch property.

Recover low-rank approximation using two-sided sampling and projection-cost-preserving sketch property.



FINAL ALGORITHM

Recover low-rank approximation using two-sided sampling and projection-cost-preserving sketch property.



- View each entry of \mathbf{A} as encoding a large amount of information about its square root \mathbf{B} . In particular $\mathbf{a}_{ij} = \mathbf{b}_i^T \mathbf{b}_j$.

SUMMARY OF MAIN IDEAS

- View each entry of \mathbf{A} as encoding a large amount of information about its square root \mathbf{B} . In particular $\mathbf{a}_{ij} = \mathbf{b}_i^T \mathbf{b}_j$.
- Use this view to find a low-rank approximation to \mathbf{B} using sublinear accesses to \mathbf{A} .

SUMMARY OF MAIN IDEAS

- View each entry of \mathbf{A} as encoding a large amount of information about its square root \mathbf{B} . In particular $\mathbf{a}_{ij} = \mathbf{b}_i^T \mathbf{b}_j$.
- Use this view to find a low-rank approximation to \mathbf{B} using sublinear accesses to \mathbf{A} .
- Since \mathbf{B} has the same singular vectors as \mathbf{A} and $\sigma_i(\mathbf{B}) = \sqrt{\sigma_i(\mathbf{A})}$, a low-rank approximation of \mathbf{B} can be used to find one for \mathbf{A} , albeit with a \sqrt{n} factor loss in quality.

SUMMARY OF MAIN IDEAS

- View each entry of \mathbf{A} as encoding a large amount of information about its square root \mathbf{B} . In particular $\mathbf{a}_{ij} = \mathbf{b}_i^T \mathbf{b}_j$.
- Use this view to find a low-rank approximation to \mathbf{B} using sublinear accesses to \mathbf{A} .
- Since \mathbf{B} has the same singular vectors as \mathbf{A} and $\sigma_i(\mathbf{B}) = \sqrt{\sigma_i(\mathbf{A})}$, a low-rank approximation of \mathbf{B} can be used to find one for \mathbf{A} , albeit with a \sqrt{n} factor loss in quality.
- Obtain near-optimal complexity using ridge leverage scores to sample both rows and columns of \mathbf{A} .


OPEN QUESTIONS

- What else can be done for PSD matrices? We give applications to ridge regression, but what other linear algebraic problems require a second look?


- What else can be done for PSD matrices? We give applications to ridge regression, but what other linear algebraic problems require a second look?
- Are there other natural classes of matrices that admit sublinear time low-rank approximation?

OPEN QUESTIONS

- What else can be done for PSD matrices? We give applications to ridge regression, but what other linear algebraic problems require a second look?
- Are there other natural classes of matrices that admit sublinear time low-rank approximation?
 - Starting points are matrices that break the $\Omega(\text{nnz}(\mathbf{A}))$ time lower bound: e.g. binary matrices, diagonally dominant matrices.



- What else can be done for PSD matrices? We give applications to ridge regression, but what other linear algebraic problems require a second look?
- Are there other natural classes of matrices that admit sublinear time low-rank approximation?
 - Starting points are matrices that break the $\Omega(\text{nnz}(\mathbf{A}))$ time lower bound: e.g. binary matrices, diagonally dominant matrices.



- What can we do when we have PSD matrices with additional structure? E.g. kernel matrices.

Thanks! Questions?